

Height of an AVL tree

Define $N(h)$ to be the *least* number of nodes in a height- h AVL tree.

One subtree must have height at least $h - 1$, the other at least $h - 2$:

$$N(h) = \begin{cases} 1 + N(h-1) + N(h-2), & h \geq 1 \\ 1, & h = 0 \\ 0, & h = -1 \end{cases}$$

What sequence does this look like? The Fibonacci sequence!

$$N(h) = F_{h+3} - 1 = \left\lceil \frac{\varphi^{h+3}}{\sqrt{5}} \right\rceil - 1, \text{ where } \varphi = \frac{1 + \sqrt{5}}{2}$$

AVL Tree Analysis

Easier lower bound on $N(h)$:

$$N(h) > 2N(h-2) > 4N(h-4) > 8N(h-6) > \dots > 2^i N(h-2i) \geq 2^{\lfloor h/2 \rfloor}$$

AVL Tree Analysis

Easier lower bound on $N(h)$:

$$N(h) > 2N(h-2) > 4N(h-4) > 8N(h-6) > \dots > 2^i N(h-2i) \geq 2^{\lfloor h/2 \rfloor}$$

Since $n > 2^{\lfloor h/2 \rfloor}$, $h \leq 2 \lg n$,

and thus an AVL tree with n nodes has height $O(\log n)$.

Also, $n \leq 2^{h+1} - 1$, so the height is $\Theta(\log n)$.

⇒ *search, insert, delete* all cost $\Theta(\log n)$.

2-3 Trees

A 2-3 Tree is like a BST with additional structural properties:

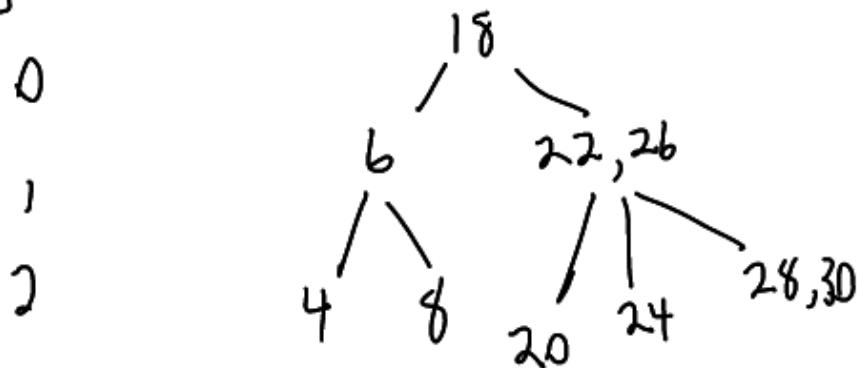
- Every node either contains *one KVP* and *two children*,
or *two KVPs* and *three children*.
- All the leaves are at the same level.
(A leaf is a node with empty children.)

Searching through a 1-node is just like in a BST.

For a 2-node, we must examine both keys and follow the appropriate path.

B-trees

level 0



An (a,b) B-tree

1 - An "ordered tree"

2 - Each internal node has at least a and at most b children

Except root which has at least 2 and at most b children

3. a node with k children has k-1 keys

4. All leaves have same level

A B-tree of 'order M' $\Rightarrow \left(\left[\frac{M}{2}\right], M\right)$

B-tree \rightarrow insert:

1 \rightarrow insert at a leaf

2 \rightarrow For overfilled nodes \Rightarrow send middle key to the parent
split to two separate nodes

B-tree Delete(g)

1-like BSTs, replace g w/ its successor / predecessor if required

2-if a node becomes underloaded

↳ if ∃ a direct sibling with extra key (more than 'a' keys) take a key from parent
parent gets a key from sibling

(a, b)-tree → at least a children / node
most b

each node → at least a-1 KVPs.
at most b-1

if $a = \lceil \frac{M}{2} \rceil$ then we have B-tree of order - M
 $b = M$

If $M=3$ $\begin{cases} a=2 \\ b=3 \end{cases}$ → 2-3 tree

Insertion in a 2-3 tree

First, we search to find the leaf where the new key belongs.

If the leaf has only 1 KVP, just add the new one to make a 2-node.

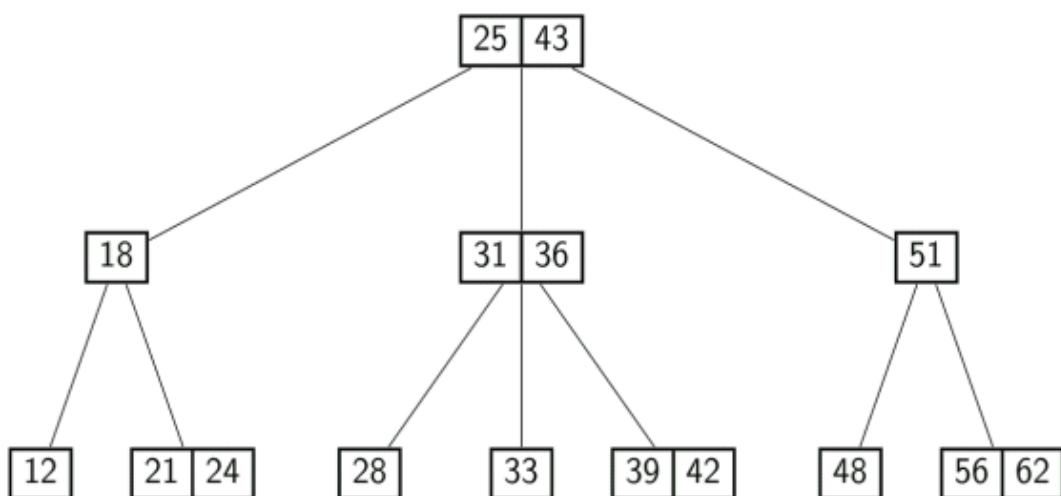
Otherwise, order the three keys as $a < b < c$.

Split the leaf into two 1-nodes, containing a and c ,
and (recursively) insert b into the parent along with the new link.

at most $b-1$ keys per node to locate the right
path to follow $\rightarrow \log b \rightarrow$ insertion takes $O(h \log b)$
(binary search)

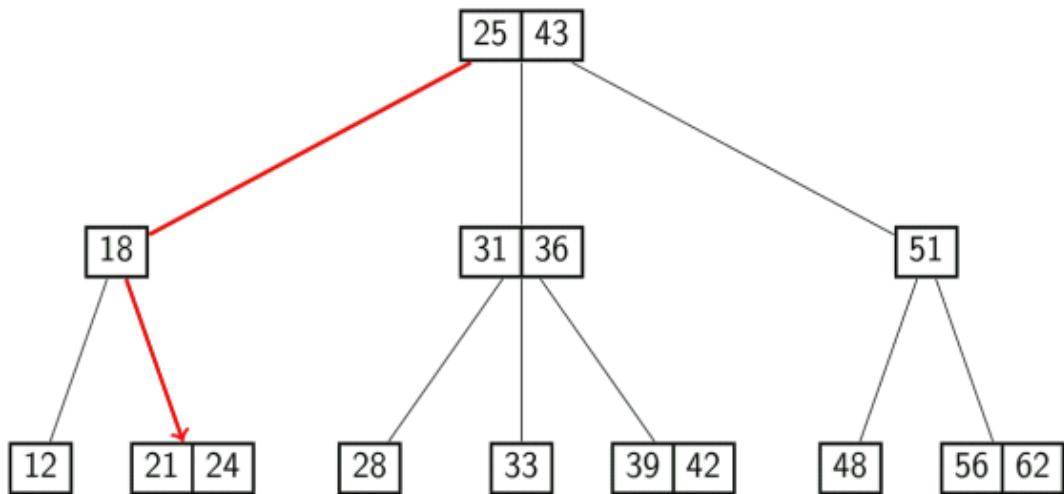
2-3 Tree Insertion

Example: `insert(19)`



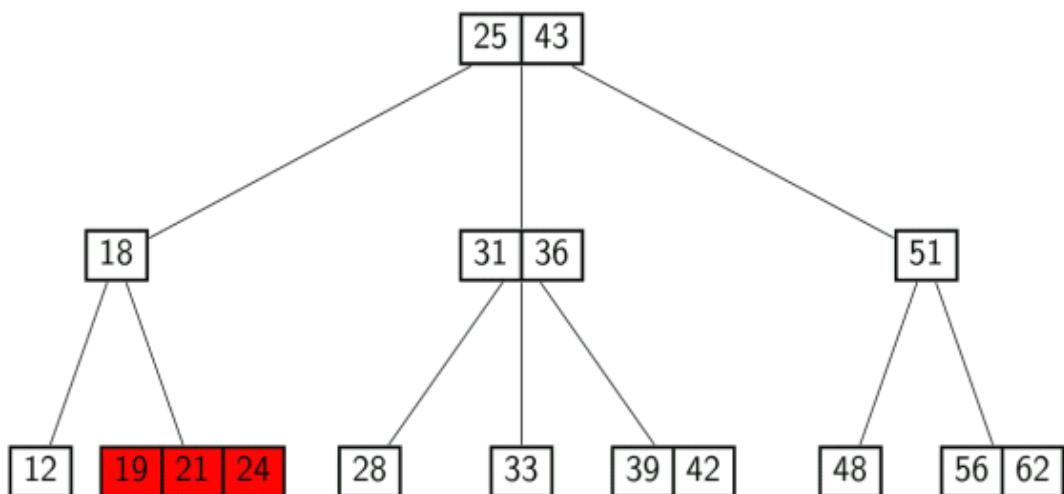
2-3 Tree Insertion

Example: *insert*(19)



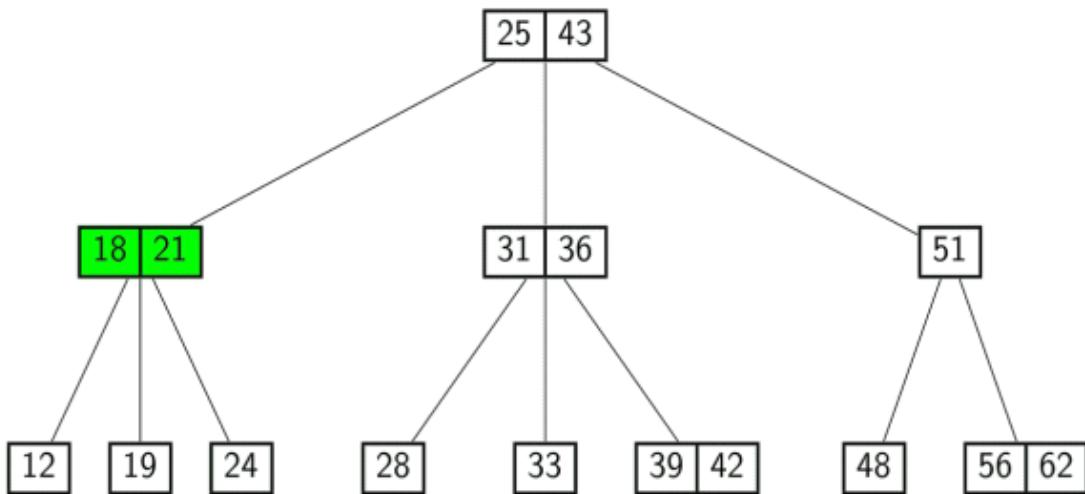
2-3 Tree Insertion

Example: *insert*(19)



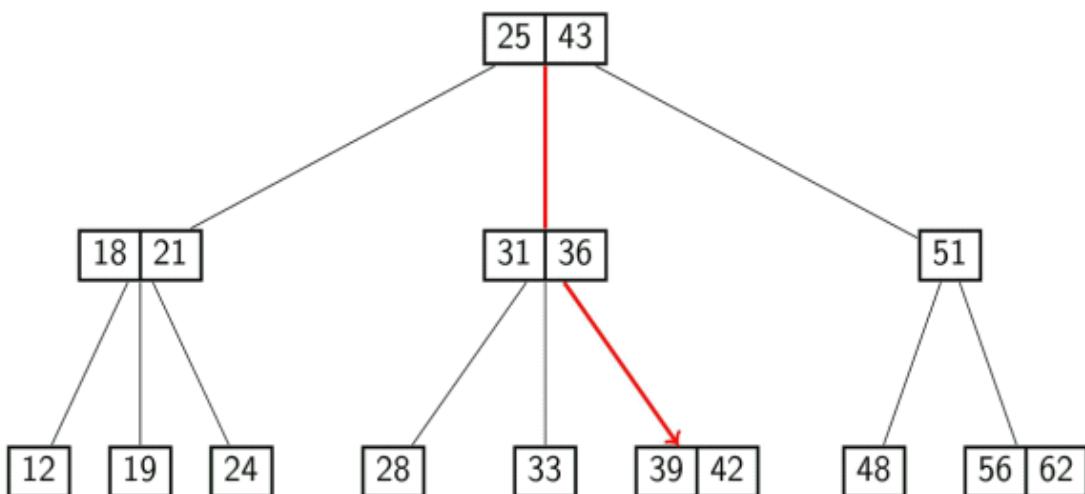
2-3 Tree Insertion

Example: *insert*(19)



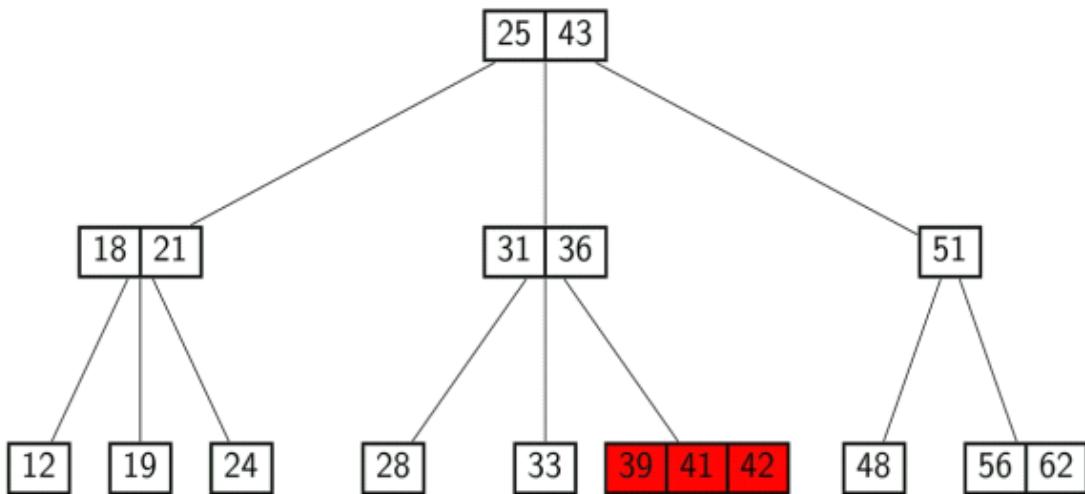
2-3 Tree Insertion

Example: *insert*(41)



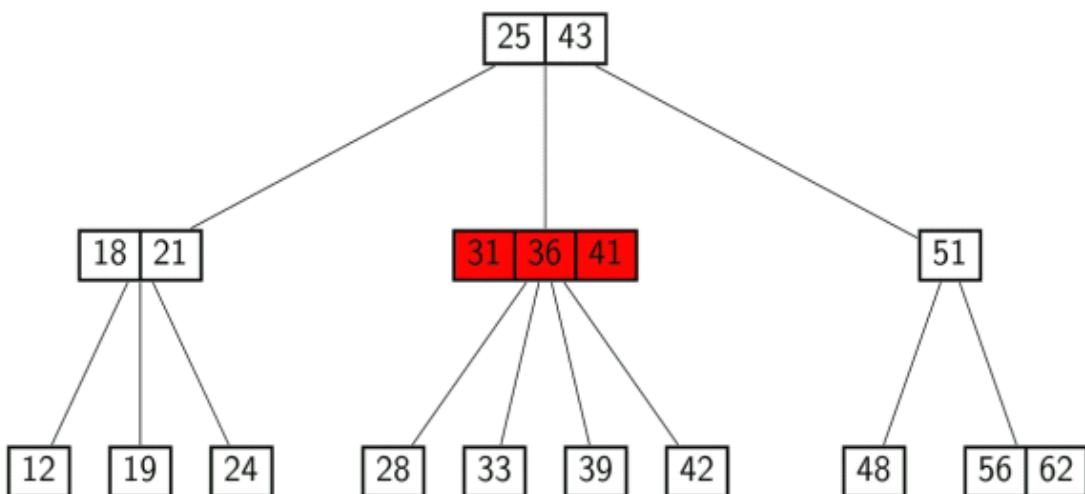
2-3 Tree Insertion

Example: *insert*(41)



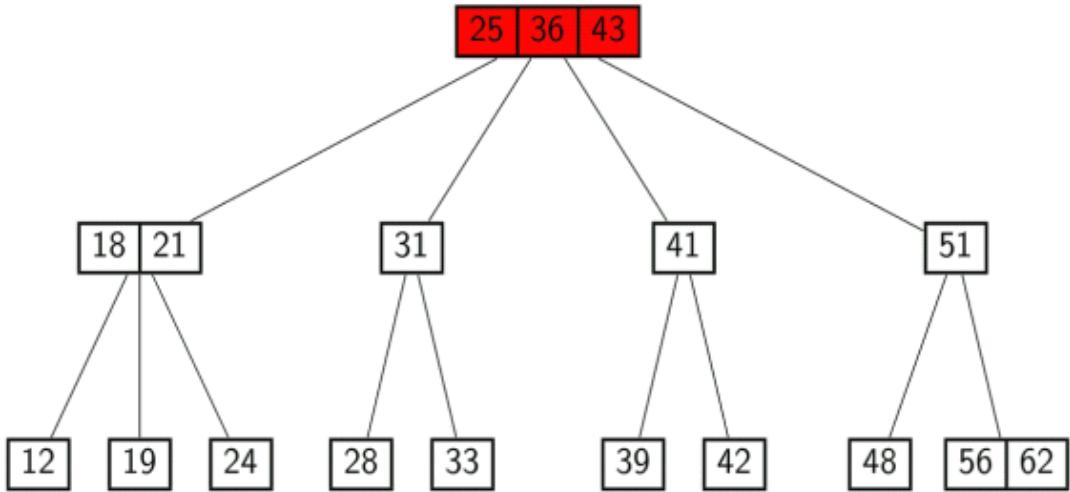
2-3 Tree Insertion

Example: *insert*(41)



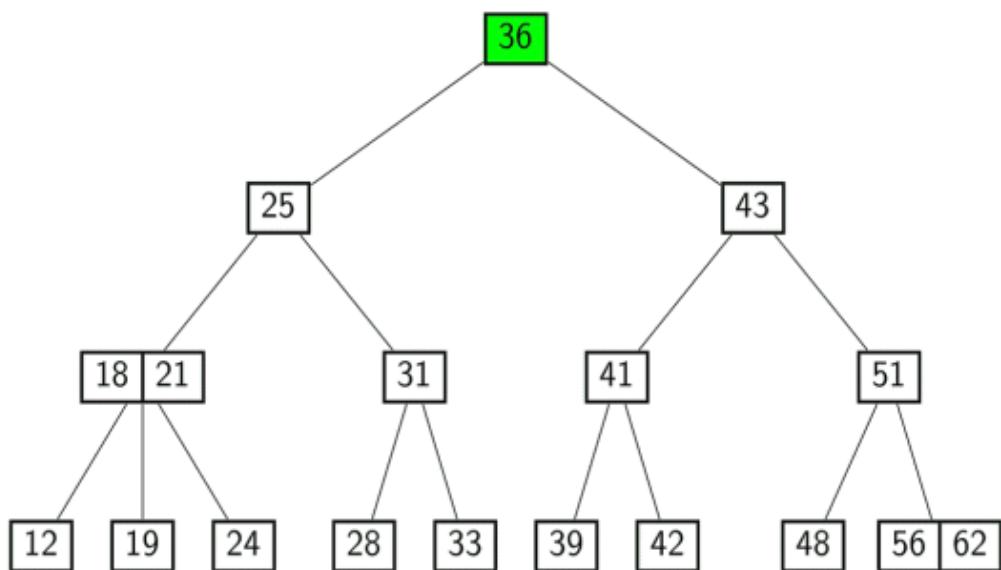
2-3 Tree Insertion

Example: *insert*(41)



2-3 Tree Insertion

Example: *insert*(41)



Deletion from a 2-3 Tree

As with BSTs and AVL trees, we first swap the KVP with its successor, so that we always delete from a leaf.

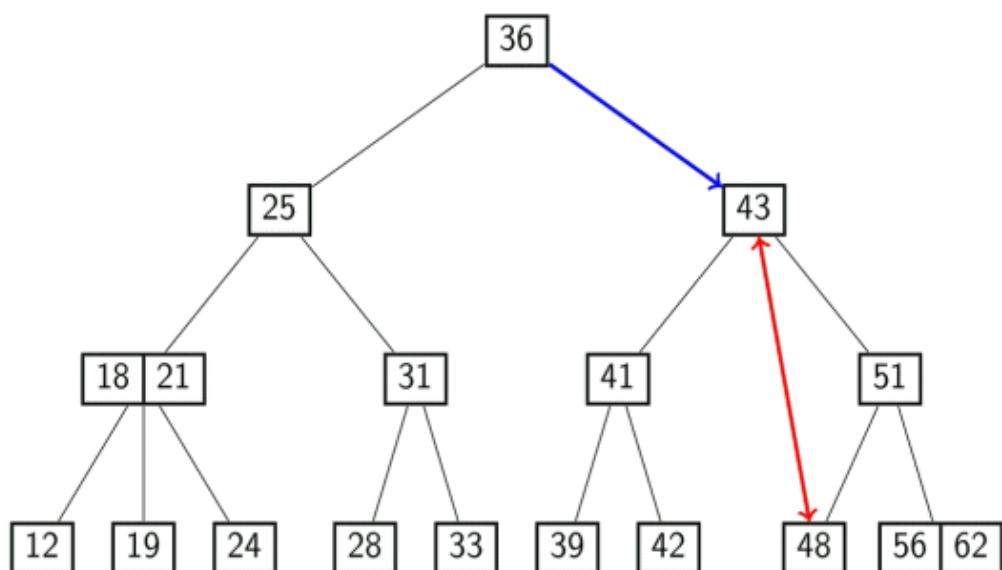
Say we're deleting KVP x from a node V :

- If V is a 2-node, just delete x .
- Elseif V has a 2-node *immediate* sibling U , perform a *transfer*:
Put the “intermediate” KVP in the parent between V and U into V , and replace it with the adjacent KVP from U .
- Otherwise, we *merge* V and a 1-node sibling U :
Remove V and (recursively) delete the “intermediate” KVP from the parent, adding it to U .

Deletion/Search/Insert $\rightarrow \Theta(h \log b)$

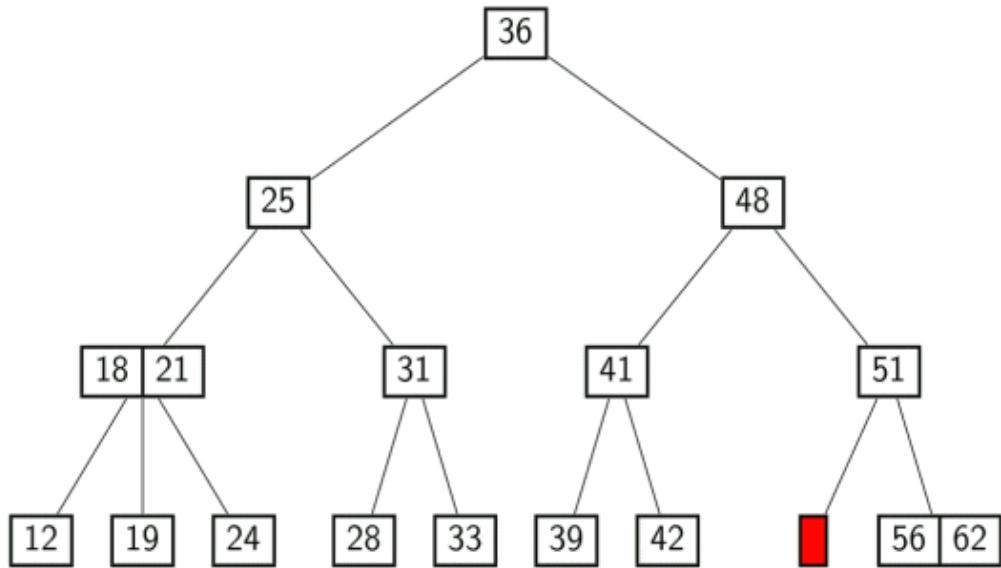
2-3 Tree Deletion

Example: *delete*(43)



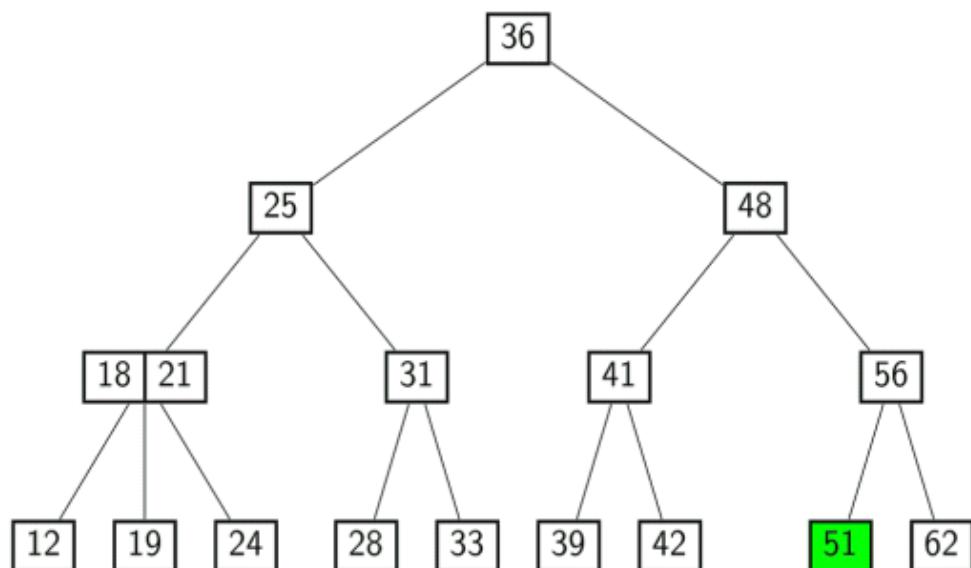
2-3 Tree Deletion

Example: `delete(43)`



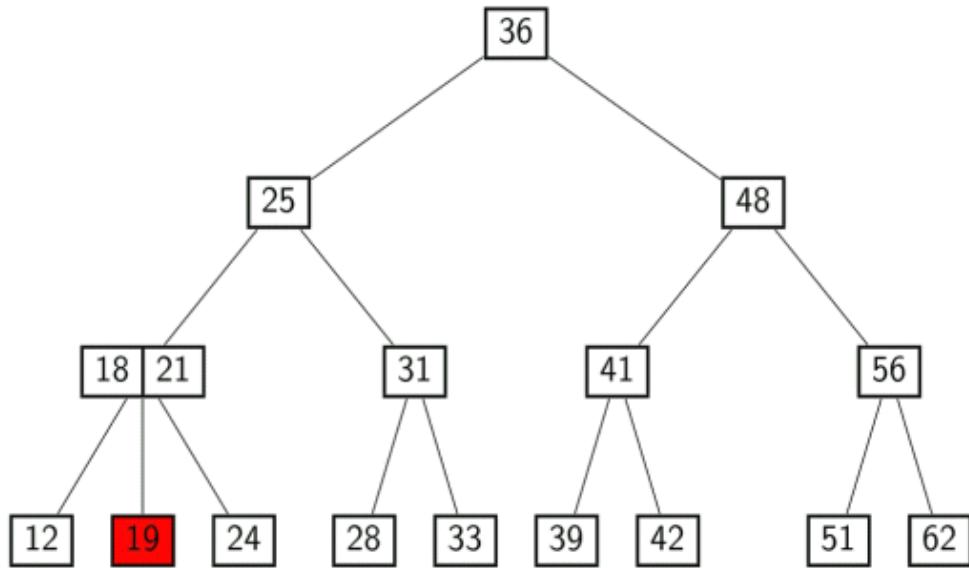
2-3 Tree Deletion

Example: `delete(43)`



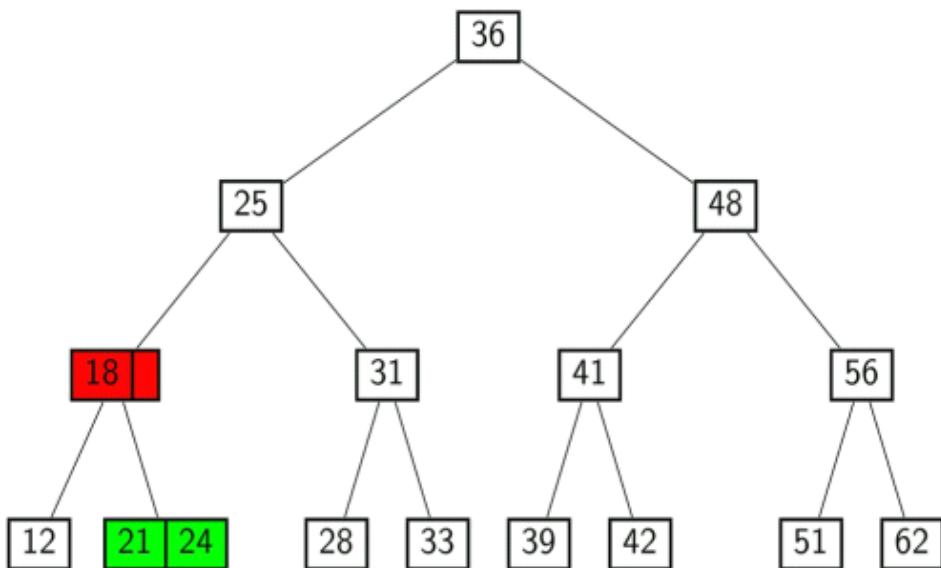
2-3 Tree Deletion

Example: `delete(19)`



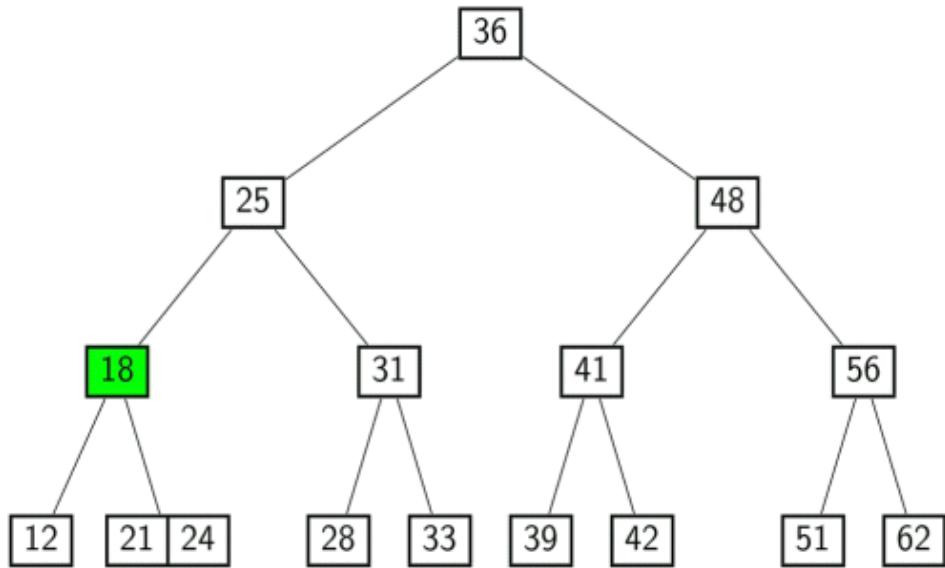
2-3 Tree Deletion

Example: `delete(19)`



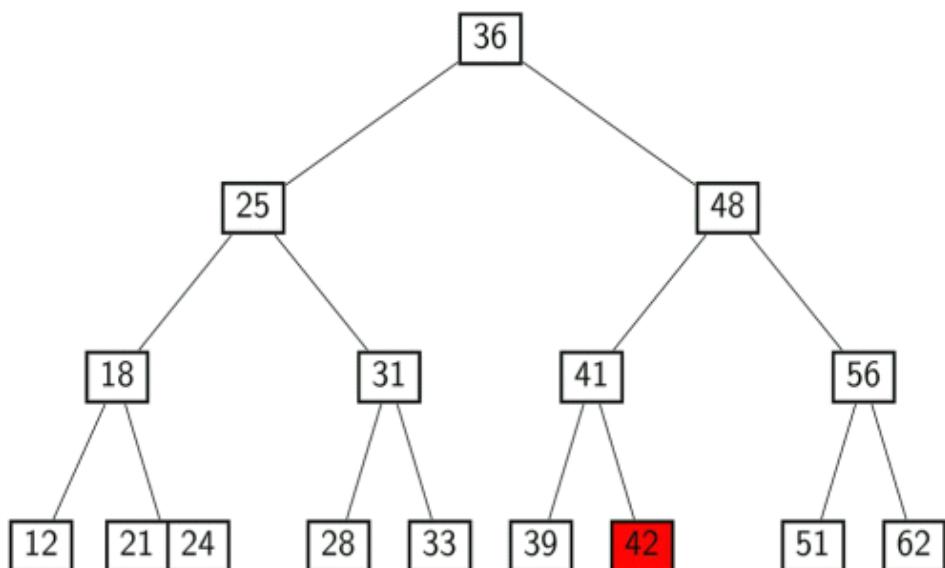
2-3 Tree Deletion

Example: `delete(19)`



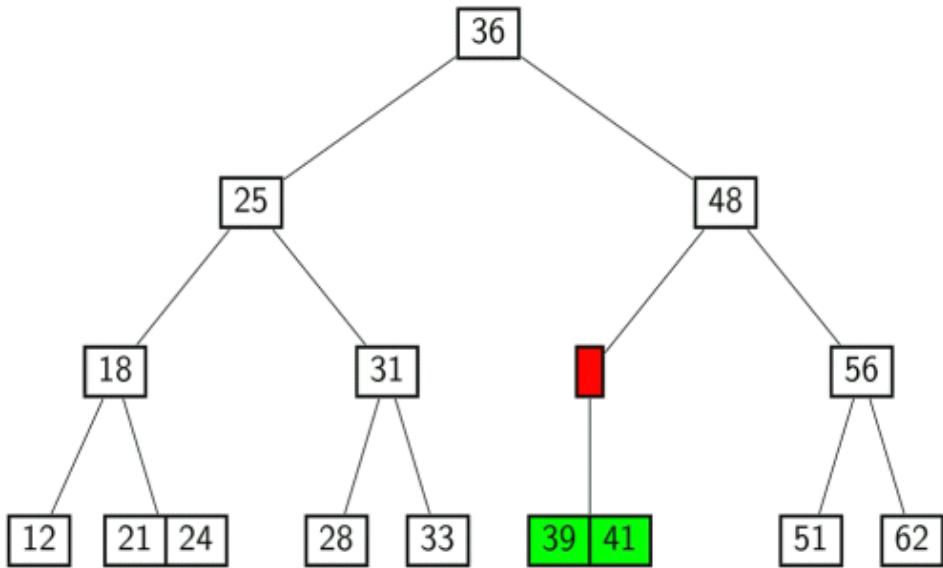
2-3 Tree Deletion

Example: `delete(42)`



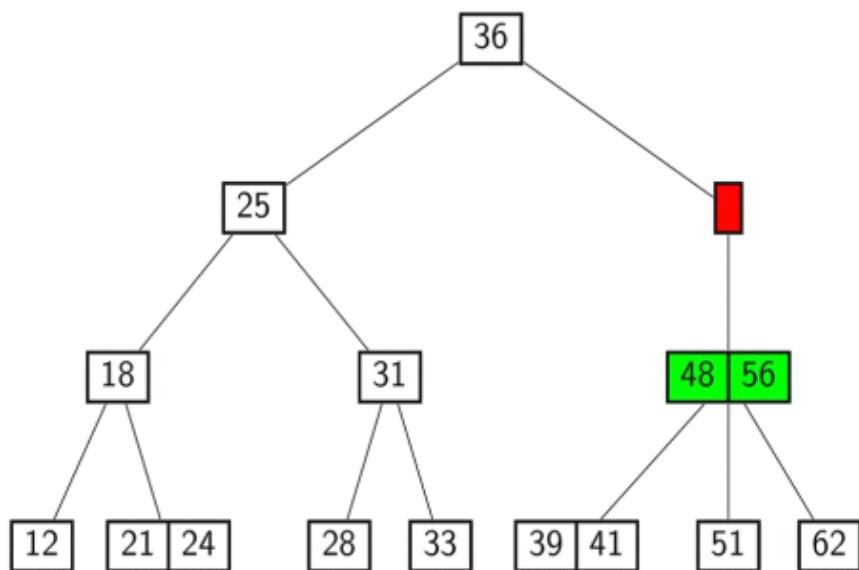
2-3 Tree Deletion

Example: `delete(42)`



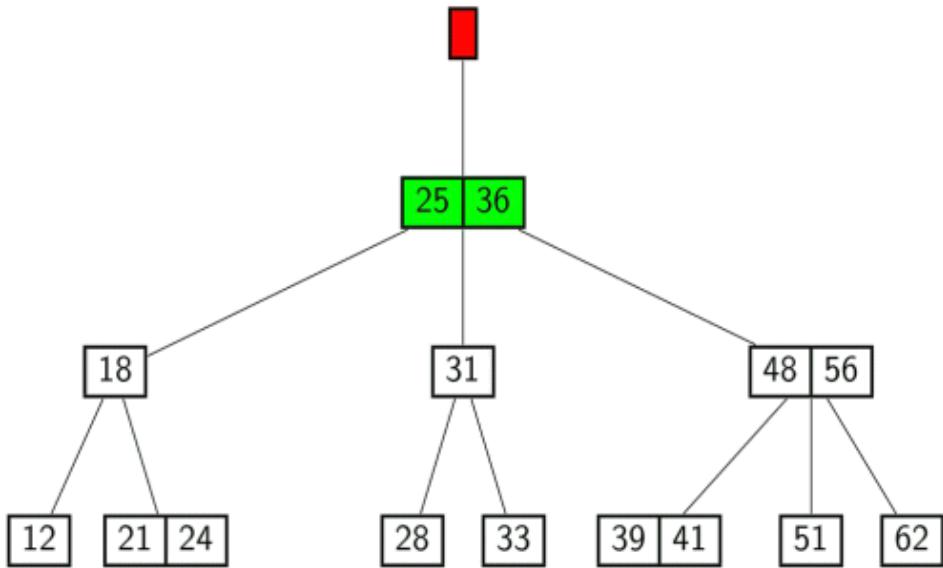
2-3 Tree Deletion

Example: `delete(42)`



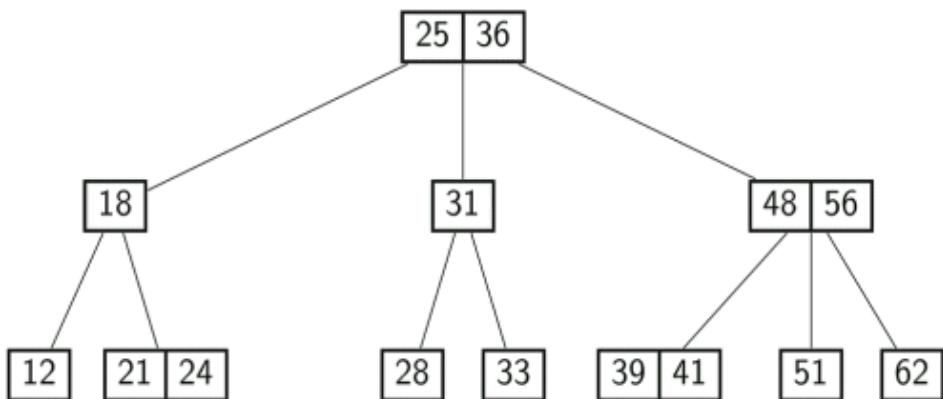
2-3 Tree Deletion

Example: `delete(42)`



2-3 Tree Deletion

Example: `delete(42)`



B-Trees

The 2-3 Tree is a specific type of B-tree:

A *B-tree of order M* is a search tree satisfying:

- Each node has between 2 and M links to (possibly empty) subtrees.
- If a node has k links, then it stores $k - 1$ key-value pairs (KVPs).
- All nodes that store links to empty subtrees are on the same level.
- Each non-root node has links to at least $\lceil M/2 \rceil$ subtrees.
- The root has links to at least 2 subtrees.

Some people call this a B-tree of *min-size* $\lceil M/2 \rceil$, or a $(\lceil M/2 \rceil, M)$ -tree.
A 2-3 tree has $M = 3$.

search, insert, delete work just like for 2-3 trees.

Height of a B-tree

What is the least number of KVPs in a height- h B-tree?

Level	Nodes	Links/node	KVP/node	KVPs on level
0	1	2	1	1
1	2	$M/2$	$M/2 - 1$	$2(M/2 - 1)$
2	$2(M/2)$	$M/2$	$M/2 - 1$	$2(M/2)(M/2 - 1)$
3	$2(M/2)^2$	$M/2$	$M/2 - 1$	$2(M/2)^2(M/2 - 1)$
...
h	$2(M/2)^{h-1}$	$M/2$	$M/2 - 1$	$2(M/2)^{h-1}(M/2 - 1)$

$$\text{Total: } n \geq 1 + 2 \sum_{i=0}^{h-1} (M/2)^i (M/2 - 1) = 2(M/2)^h - 1$$

Therefore height of tree with n nodes is $\Theta((\log n)/(\log M))$.

Analysis of B-tree operations

Assume each node stores its KVPs and child-pointers in a dictionary that supports $O(\log M)$ search, insert, and delete.

Then *search*, *insert*, and *delete* work just like for 2-3 trees, and each require $\Theta(\text{height})$ node operations.

Total cost is $O\left(\frac{\log n}{\log M} \cdot (\log M)\right) = O(\log n)$.

$$h = \frac{\log n}{\log M} \quad b = M$$

(order M B-tree)

Dictionaries in external memory

Tree-based data structures have poor *memory locality*:

If an operation accesses m nodes, then it must access m spaced-out memory locations.

Observation: Accessing a single location in *external memory* (e.g. hard disk) automatically loads a whole block (or “page”).

In an AVL tree or 2-3 tree, $\Theta(\log n)$ pages are loaded in the worst case.

If M is small enough so an M -node fits into a single page, then a B-tree of order M only loads $\Theta((\log n)/(\log M))$ pages.

This can result in a *huge* savings:
memory access is often the largest time cost in a computation.

B-tree variations

Max size M : Permitting one additional KVP in each node allows *insert* and *delete* to avoid *backtracking* via *pre-emptive splitting* and *pre-emptive merging*.

Red-black trees: Identical to a B-tree with minsize 1 and maxsize 3, but each 2-node or 3-node is represented by 2 or 3 binary nodes, and each node holds a “color” value of red or black.

B⁺-trees: All KVPs are stored at the leaves (interior nodes just have keys), and the leaves are linked sequentially.

RAM Machine \rightarrow preliminary operations take constant time
 \hookrightarrow if all data on RAM, then AVL & B-tree have the same complexity $O(\log n)$

\hookrightarrow if data is huge \Rightarrow on disk \Rightarrow you need to have disk access

\hookrightarrow if we count # of disk access

\hookrightarrow AVL tree. $\Theta(\log(n))$ disk access in worst case

\hookrightarrow B-tree, assuming M keys fit in a block

$\hookrightarrow \Theta(h) \in \Theta\left(\frac{\log n}{\log m}\right)$ disk accesses per operation

Red-Black tree

↳ BST Realization of (2,4) B.trees

↳ search / insert / delete $\Theta(\log n)$

↳ less structured than BSTs

Search intensive \rightarrow AVL is better

insert / delete \rightarrow RB tree is better